

DeepSeek V4.1 Flash

Step-by-step setup guide (official DeepSeek API)

Sign-up / Console
platform.deepseek.com
API Base URL (OpenAI format)
https://api.deepseek.com
API Base URL (Anthropic format)
https://api.deepseek.com/anthropic
Model Name
deepseek-flash
Model Scale
552B total parameters (MoE) 1M token context up to 384K output tokens

Important: Many third-party sites report new accounts get a 5M free-token grant (no card needed), but DeepSeek's own current pricing page does not explicitly guarantee this for every account — grants can be account-, region-, or campaign-specific. Check your Billing/Balance page in the console right after signup to see your actual granted balance before relying on it.

Step 1 — Create your account

Go to platform.deepseek.com, click Sign Up, and register with your email. Verify your email address to activate the account.

Step 2 — Check your granted balance

Log in to the developer console and open the Billing or Balance section. Note whatever granted balance appears there (if any), along with its expiry — this tells you exactly what free usage you actually have, rather than relying on third-party estimates.

Step 3 — Create your API key

In the console, go to API Keys, create a new key, name it (e.g. “development”), and copy it immediately — it will not be shown again.

Step 4 — Make your first API call

Use the model name `deepseek-flash` — this always points to the current DeepSeek-V4.1-Flash release. The API is OpenAI/Anthropic-compatible, so existing SDKs work by just changing the base URL and key.

```
curl https://api.deepseek.com/chat/completions \
-H "Content-Type: application/json" \
-H "Authorization: Bearer YOUR_API_KEY" \
-d '{
  "model": "deepseek-flash",
  "messages": [
    {"role": "system", "content": "You are a helpful assistant."},
    {"role": "user", "content": "Hello!"}
  ],
  "thinking": {"type": "enabled"},
  "reasoning_effort": "high",
  "stream": false
}'
```

Step 5 — Toggle thinking mode (optional)

The model supports both non-thinking and thinking modes (thinking is default). Set `"thinking": {"type": "enabled"}` or `"disabled"` depending on whether you want step-by-step reasoning shown. Use `reasoning_effort` (up to `"xhigh"`) to control how much reasoning the model applies before answering.

Step 6 — Use it inside coding agent tools (optional)

DeepSeek's API is directly supported by popular agent tools like Claude Code, GitHub Copilot, and OpenCode — point them at the DeepSeek base URL and API key and select `deepseek-flash` as the model, no custom code required. See DeepSeek's Agent Integrations guide for tool-specific setup.

Step 7 — Watch peak vs off-peak pricing

DeepSeek bills half-price during off-peak hours (most of the week) and full price during peak hours (01:00–04:00 and 06:00–10:00 UTC, Monday–Friday). Scheduling non-urgent work in off-peak windows stretches any free or paid balance further.

Step 8 — Monitor usage and top up if needed

Once any granted balance is used up, requests return an HTTP 402 Insufficient Balance error. Check your usage dashboard regularly, and top up via card or PayPal on the platform if you want to continue past the free allowance.

Notes: DeepSeek-V4.1-Flash is DeepSeek's newest sparse Mixture-of-Experts model (552B total parameters, 1M token context), which DeepSeek states surpasses their own previous V4 Pro flagship on performance, speed, and task completion time — this claim comes from DeepSeek's own release notes and has not yet been independently verified. Free-tier credit availability varies by account and is not explicitly guaranteed on DeepSeek's current public pricing page — always confirm your actual balance in the console before relying on it. Pricing and model behavior can change; check platform.deepseek.com and api-docs.deepseek.com for the latest details.