

Jev by TypeSafe AI, via Vercel

Step-by-step setup guide (free through Vercel AI Gateway)

Sign-up / Console

vercel.com → Dashboard → AI Gateway

Model ID

typesafe-ai/jev

Promotional Pricing Ends

September 25, 2026

Standard Pricing (after Sept 25)

\$0.04 per 1M input tokens | Output: Free

Note: Jev is not a chatbot — it does not generate text. It returns typed decisions (Choice, Score, or Boolean) with calibrated confidence scores, meant for use inside application logic (routing, classification, verification), not for conversational replies.

Step 1 — Create a Vercel account

Go to vercel.com and sign up or log in. A free account is enough to access the AI Gateway — no card required for this.

Step 2 — Open AI Gateway

From your dashboard, navigate to **AI Gateway**, then click “**Get API key**” (found on the Jev model page or the AI Gateway Quick Start).

Step 3 — Generate and copy your API key

Create a new Gateway API key and copy it immediately. This single key works across every model on Vercel's AI Gateway, not just Jev.

Step 4 — Install the AI SDK

In your project, install Vercel's AI SDK, which includes the experimental evaluation API used to call Jev.

```
npm install ai
```

Step 5 — Make your first call

Use the `evaluate` function with model ID `typesafe-ai/jev`. Supply a “state” (the situation/text) and one or more typed “questions.”

```
import { experimental_evaluate as evaluate } from 'ai';

const result = await evaluate({
  model: 'typesafe-ai/jev',
  state: 'The support agent issued a full refund to the customer.',
  questions: {
    refunded: {
      type: 'boolean',
      instructions: 'Was a refund issued?',
    },
  },
});

console.log(result.answers);
// { refunded: { type: 'boolean', probability: 0.99 } }
```

Step 6 — Set your API key as an environment variable

Store your key as `AI_GATEWAY_API_KEY` in your project's environment variables (e.g. `.env` file). The AI SDK picks it up automatically — no extra configuration needed.

Step 7 — Use different question types

Jev supports three primitives: `boolean` (yes/no with probability), `choice` (pick one option from a defined set, with per-option probabilities), and `score` (rate against an ordered scale). You can evaluate multiple questions in parallel within a single call.

Step 8 — Know the pricing after the free window

Promotional free pricing on Vercel ends September 25, 2026. After that, Jev costs \$0.04 per million input tokens with free output — still dramatically cheaper than a standard LLM for this kind of structured-decision task.

Notes: Jev is TypeSafe AI's “System One” model, built by former OpenAI researcher Diogo Almeida and team, announced September 15, 2026 with \$40M in seed funding. TypeSafe reports Jev is 40–200x faster and up to 444x cheaper than frontier LLMs on structured-decision workflows — these are the company's own launch benchmarks and have not yet been independently verified. Jev does not generate prose; pair it with a generative model when a written response is needed. Always confirm current pricing and terms at vercel.com/ai-gateway/models/jev and typesafe.ai before relying on it in production.